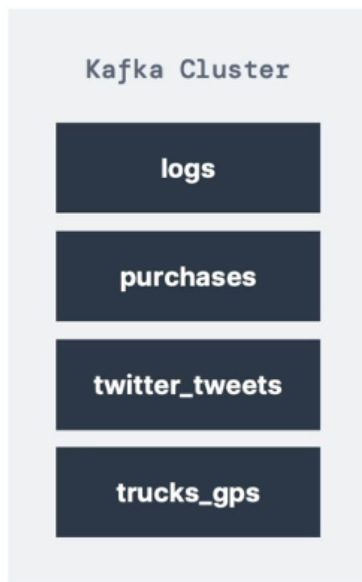


Kafka Theory Part 1

▼ 1 Topics, Partitions, Offsets

- ▼ Kafka Topic: a particular stream of data
 - Lika a table in DB without constraints



- Identified by name
- Any kind of message format
- sequence of messages is called data stream
- Cannot query Topic
- Immutable

▼ Partitions and Offsets

- Topics are split into partitions
- Messages in partitions are ordered




- Each message within a partition gets an inc id called offset
 - Offset only have a meaning for a specific partition
- ▼ Example: truck_gps - GPS pos
- 10 partitions

- Truck send msg every 20 sec
- Location service & notification service consumer

▼ 2 Producers and Message Keys

▼ Producer

- Producers write data to topics



- Producer know to which partition to write to
- Producer auto recover if broker failures

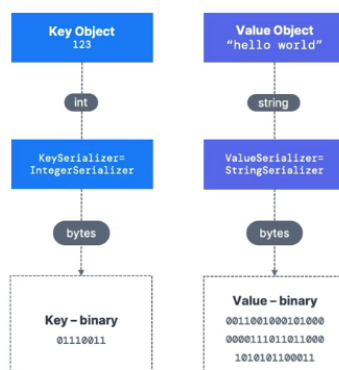
▼ Message keys

- Producers can choose to send a key with msg
- if key=null, round robin
- if key !=null all message will go to same partition



▼ Message

- Message Anatomy: <Key-binary, Value-binary>, CompressionType, Headers, Partition+offset, TimeStamp
- Kafka Message Serializer: kafka accepts only bytes as inputs and sends bytes. Message Seriazlization means transforming obj to bytes



- Common serializer: string (json), int, float, avro, protobuf

- Kafka Key hashing



- **Key Hashing** is the process of determining the mapping of a key to a partition
- In the default Kafka partitioner, the keys are hashed using the **murmur2 algorithm**, with the formula below for the curious:


```
targetPartition = Math.abs(Utils.murmur2(keyBytes)) % (numPartitions - 1)
```

▼ 4 Consumer Groups and Consumer Offsets

- All the consumers in app read data as consumer groups
- if consumers > partition, some consumers will be inactive
- Multiple consumer group per Topic is accepted
- Kafka stores the offsets at which a consumer group has been reading
- Offsets committed are in Kafka topic named `__consumer_offsets`
- ▼ When a consumer consumers it should periodically commit the offsets
 - If the consumer dies, it can pick up from where it left
- ▼ Delivery semantics for consumers
 - Default: auto commit offsets
 - ▼ Manual Commit
 - Atleast once: offsets commit after processing. if goes wrong again read. So need idempotent(processing again wont impact)
 - At most once: offsets are committed as soon as messages are received. if processing goes wrong some msg will be lost
 - Exactly Once: kafka => kafka workflows: use Transactional API

▼ 3 Consumers and Deserialization

- Consumers read data from a topic - pull model



- Consumer know which partition to read from
- In case of broker failure consumer knows to recover
- Data is read in order within each partition

- Consumer Deserializer - deserialize from binary to object

